

# KQL

Kusto Query Language

## Complete Cheat Sheet

Beginner → Intermediate → Advanced → Master

*Azure Data Explorer • Log Analytics • Microsoft Sentinel*

*Covers: Filtering • Aggregations • Joins • Time Series • Threat Detection • Analytics Rules*

---

# 1. KQL FUNDAMENTALS

## The Pipeline Mental Model

KQL works like a pipeline. Data flows left-to-right through | (pipe) operators. Each step filters or transforms the data.

```
// Table → filter → transform → output
SignIns
| where TimeGenerated > ago(1h)
| project UserPrincipalName, IPAddress, ResultType
| sort by TimeGenerated desc
```

■ Always put your TimeGenerated filter FIRST — it limits data scanned and makes queries much faster.

## Time Filters — ago()

```
// Positive time - most common in log analysis
| where TimeGenerated > ago(1h) // last 1 hour
| where TimeGenerated > ago(24h) // last 24 hours
| where TimeGenerated > ago(7d) // last 7 days
| where TimeGenerated > ago(30m) // last 30 minutes
// Absolute time range
| where TimeGenerated between (datetime(2024-01-01) .. datetime(2024-01-31))
```

## Filtering with where

```
// Exact filter
| where ResultType == '0'
// Not equal
| where ResultType != '0'
// Multiple conditions / AND
| where ResultType != '0' and Location == 'Russia'
// OR condition
| where Location == 'Russia' or Location == 'Nigeria'
// Match a list of values
| where Location in ('Russia', 'Nigeria', 'China', 'Iran')
// Excluded values
| where Location !in ('United States', 'United Kingdom')
// Containing text (allow ~)
| where UserPrincipalName contains 'admin'
// Has whole word (Factor = prefer this)
| where UserPrincipalName has 'admin'
// Starts with / ends with
| where UserPrincipalName startswith 'svc'
// Matching terms - has_and / has_all
| where AppDisplayName has_any ('Teams', 'SharePoint', 'Portal')
```

## Selecting Columns — project

```
// Pick only the columns you need
| project TimeGenerated, UserPrincipalName, IPAddress, ResultType
// Remove columns while projecting
| project Time=TimeGenerated, User=UserPrincipalName, IP=IPAddress
// Keep all columns EXCEPT some
| project-away AttackPattern, IsInteractive
```

```
// add a new calculated column
| extend IsFailed = iff(ResultType != '0', true, false)
```

## Sorting and Limiting

```
| sort by TimeGenerated desc // newest first
| sort by TimeGenerated asc // oldest first
| top 10 by TimeGenerated desc // top 10 most recent
| take 50 // any 50 rows (no order)
```

## Essential Operators Quick Reference

Operator	Description	Example
where	Filter rows	where EventID == 4625
project	Select columns	project User, IP
extend	Add new column	extend IsFail = ResultType != '0'
sort / order by	Sort results	sort by count_desc
top	Top N rows	top 10 by count_
take / limit	Return N rows	take 100
count	Count all rows	count
distinct	Unique values	distinct UserPrincipalName
summarize	Aggregate / group by	summarize count() by User
render	Visualise as chart	render timechart

● INTERMEDIATE

## 2. AGGREGATIONS & SUMMARIZE

### summarize — The Most Important KQL Operator

```
// Basic count
SignIns | summarize count()

// Count per group (like .dot dcount dv)
SignIns | summarize EventCount = count() by UserPrincipalName

// More complex aggregations at once
SignIns
| summarize
    TotalLogins = count(),
    FailedLogins = countif(ResultType != '0'),
    SuccessLogins = countif(ResultType == '0'),
    UniqueIPs = dcount(IPAddress)
    by UserPrincipalName
| sort by TotalLogins desc
```

### Aggregation Functions

Function	Description	Example
count()	Total row count	summarize count()
countif(cond)	Count rows matching condition	countif(ResultType != '0')
dcount(col)	Distinct / unique count	dcount(IPAddress)
sum(col)	Sum of a column	sum(Duration)
avg(col)	Average value	avg(ResponseTime)
min(col)	Minimum value	min(TimeGenerated)
max(col)	Maximum value	max(TimeGenerated)
make_set(col)	List of unique values	make_set(Location)
make_list(col)	List of all values	make_list(IPAddress)
stdev(col)	Standard deviation	stdev(Duration)
percentile(col,n)	Nth percentile	percentile(Latency,95)

### Time Bucketing with bin()

bin() groups timestamps into fixed time windows — essential for trend analysis.

```
// Events per hour
SignIns
| summarize Events = count() by bin(TimeGenerated, 1h)
| render timechart

// Failed vs successful per day
SignIns
| summarize
    Failed = countif(ResultType != '0'),
    Success = countif(ResultType == '0')
```

```
by bin(TimeGenerated, 1d)
| render timechart
```

## let — Variables and Reusable Logic

```
// Store values in variables
let timeframe = 24h;
let threshold = 10;

let badCountries = dynamic(['Russia','Nigeria','China','Iran','Belarus']);

SignIns
| where TimeGenerated > ago(timeframe)
| where Location in (badCountries)
| summarize Attempts = count() by UserPrincipalName, IPAddress
| where Attempts > threshold
| sort by Attempts desc
```

■ Use `let` to define thresholds and time windows at the top of your query — makes tuning analytics rules much easier.

## String Functions

Function	Description	Example
<code>tostring(x)</code>	Convert to string	<code>tostring(EventID)</code>
<code>toint(x)</code>	Convert to integer	<code>toint(ResultType)</code>
<code>tolower(x)</code>	Lowercase	<code>tolower(UserPrincipalName)</code>
<code>toupper(x)</code>	Uppercase	<code>toupper(Location)</code>
<code>strlen(x)</code>	String length	<code>strlen(CommandLine)</code>
<code>substring(x,i,n)</code>	Extract substring	<code>substring(UPN, 0, 5)</code>
<code>split(x,delim)</code>	Split string to array	<code>split(UPN, '@')</code>
<code>strcat(a,b)</code>	Concatenate strings	<code>strcat(User, ' - ', IP)</code>
<code>replace_string()</code>	Find and replace	<code>replace_string(col,'old','new')</code>
<code>extract(regex,1,x)</code>	Extract with regex	<code>extract(@'(\d+\.\d+)', 1, IP)</code>

● ADVANCED

## 3. JOINS, UNIONS & MULTI-TABLE QUERIES

### join — Combining Two Tables

```
// inner join = only rows that match in BOTH tables

SignIns
| where ResultType == '0' // successful logins
| join kind=inner (
    SignIns
    | where ResultType != '0' // failed logins
    | summarize FailCount = count() by UserPrincipalName
    | where FailCount > 5
) on UserPrincipalName
| project TimeGenerated, UserPrincipalName, IPAddress, FailCount
// Result: accounts that failed 5+ times then succeeded = compromise signal
```

### Join Types Quick Reference

Join Kind	Returns	Use Case
inner	Rows matching in both tables	Find overlap between two events
leftouter	All left rows + matching right rows	Enrich alerts with user details
rightouter	All right rows + matching left rows	Find unmatched alerts
fullouter	All rows from both tables	Complete comparison
leftanti	Left rows NOT in right table	Find logins with no prior fails
rightanti	Right rows NOT in left table	Find events with no resolution
leftsemi	Left rows that have match in right	Filter by existence

### union — Combine Multiple Tables

```
// Stack two tables together (like SQL UNION ALL)
union SignIns, SecurityEvent
| where TimeGenerated > ago(24h)
| summarize count() by Type, bin(TimeGenerated, 1h)
| render timechart

// UNION with wildcard = all tables starting with 'Security'
union Security*
| where TimeGenerated > ago(1h)
| summarize count() by Type
```

### mv-expand — Expand Arrays into Rows

```
// make set() created arrays - mv-expand unpacks them
SignIns
| summarize Countries = make_set(Location) by UserPrincipalName
| mv-expand Countries
| project UserPrincipalName, Country = tostring(Countries)
```

### parse — Extract Values from Unstructured Text

```
// Extract fields from front-text log messages
SignIns
| parse ResultDescription with 'Error: ' ErrorCode - 'ErrorMsg'
| project UserPrincipalName, ErrorCode, ErrorMsg
// Trailing regex extract
| extend Domain = extract(@'@(.+)$', 1, UserPrincipalName)
```

## iff & case — Conditional Columns

```
// iff = simple if/else
| extend Status = iff(ResultType == '0', 'Success', 'Failed')
// case = multiple conditions / like switch/case
| extend RiskLabel = case(
    RiskLevelDuringSignIn == 'high', 'CRITICAL',
    RiskLevelDuringSignIn == 'medium', 'WARNING',
    RiskLevelDuringSignIn == 'low', 'INFO',
    'NORMAL' )
```

● MASTER

## 4. SENTINEL THREAT DETECTION QUERIES

### Brute Force Detection

```
let timeframe = 1h;
let threshold = 10;

SignIns
| where TimeGenerated > ago(timeframe)
| where ResultType != '0'
| summarize FailedAttempts = count() by UserPrincipalName, IPAddress, Location
| where FailedAttempts > threshold
| sort by FailedAttempts desc
```

### Password Spray Detection

```
// One IP targeting many accounts with few attempts each
SignIns
| where TimeGenerated > ago(1h)
| where ResultType != '0'
| summarize
    AccountsTargeted = dcount(UserPrincipalName),
    TotalAttempts = count()
    by IPAddress, Location
| where AccountsTargeted > 5 and TotalAttempts < 100
| sort by AccountsTargeted desc
```

### MFA Fatigue Detection

```
// Attacker sends MFA requests hoping user approves
SignIns
| where TimeGenerated > ago(1h)
| where ResultType == '500121' // MFA denied by user
| summarize MFADenials = count() by UserPrincipalName, IPAddress
| where MFADenials > 5
| sort by MFADenials desc
```

### Impossible Travel Detection

```
// Same user logged in from 2+ countries
SignIns
| where TimeGenerated > ago(24h)
| where ResultType == '0'
| summarize
    Countries = make_set(Location),
    CountryCount = dcount(Location),
    IPs = make_set(IPAddress)
    by UserPrincipalName
| where CountryCount > 1
| sort by CountryCount desc
```

## Account Compromise — Success After Many Failures

```
let failed =  
SignIns  
| where TimeGenerated > ago(1h)  
| where ResultType != '0'  
| summarize FailCount = count() by UserPrincipalName  
| where FailCount > 5;  
  
SignIns  
| where TimeGenerated > ago(1h)  
| where ResultType == '0'  
| join kind=inner failed on UserPrincipalName  
| project TimeGenerated, UserPrincipalName, IPAddress, Location, FailCount  
| sort by FailCount desc
```

## Sign-Ins from High-Risk Countries

```
let watchlist = dynamic(['Russia','Nigeria','China','North Korea','Iran','Belarus']);  
  
SignIns  
| where TimeGenerated > ago(24h)  
| where Location in (watchlist)  
| summarize  
Attempts = count(),  
Users = make_set(UserPrincipalName)  
by Location, IPAddress  
| sort by Attempts desc
```

● MASTER

## 5. SENTINEL ANALYTICS RULES

### Anatomy of a Sentinel Analytics Rule

```
// This is the KQL query you paste into the Analytics Rule editor
let timeframe = 1h;
let threshold = 10;

SignIns
| where TimeGenerated > ago(timeframe)
| where ResultType != '0'
| summarize FailedAttempts = count() by
UserPrincipalName, IPAddress, Location
| where FailedAttempts > threshold

// In the rule wizard:
// Run every: 1 hour
// Lookup data from: Last 1 hour
// Alert threshold: Is greater than 0
// Map entities: Account = UserPrincipalName, IP = IPAddress
```

■ Map entities in the rule wizard so Sentinel can link incidents to user profiles and IP intelligence automatically.

● INTERMEDIATE

## 6. VISUALISATIONS — render

### Chart Types

Chart Type	Best Used For	KQL
timechart	Events over time / trends	render timechart
barchart	Compare categories	render barchart
columnchart	Compare values side by side	render columnchart
piechart	Show proportions	render piechart
scatterchart	Correlate two values	render scatterchart
areachart	Volume over time	render areachart

```
// Dull example - login trend chart
SignIns
| where TimeGenerated > ago(7d)
| summarize
Failed = countif(ResultType != '0'),
Success = countif(ResultType == '0')
by bin(TimeGenerated, 1d)
| render timechart with (title='Daily Login Trend')
```

● INTERMEDIATE

## 7. COMMON SENTINEL TABLES REFERENCE

Table	Contains	Key Columns
SigninLogs	Azure AD / Entra sign-ins	UserPrincipalName, IPAddress, ResultType
AuditLogs	Azure AD changes	OperationName, InitiatedBy, TargetResources
SecurityEvent	Windows Security events	EventID, Account, Computer, LogonType
CommonSecurityLog	Firewall / IDS (CEF)	DeviceAction, SourceIP, DestinationIP
Syslog	Linux system logs	SyslogMessage, Computer, Facility
AzureActivity	Azure subscription actions	OperationName, Caller, ResourceGroup
OfficeActivity	Microsoft 365 events	Operation, UserId, ClientIP
DeviceEvents	MDE endpoint events	ActionType, DeviceName, InitiatingProcess
DeviceProcessEvents	Process creation (MDE)	FileName, ProcessCommandLine, AccountName
DeviceNetworkEvents	Network connections (MDE)	RemoteIP, RemoteUrl, ActionType
SecurityAlert	All Defender/Sentinel alerts	AlertName, AlertSeverity, Entities
ThreatIntelIndicator	Threat intel IOCs	IndicatorType, ThreatType, NetworkIP
AADRiskyUsers	Risky user detections	UserPrincipalName, RiskLevel, RiskDetail
IdentityLogonEvents	Identity-based logons (M365D)	AccountUpn, IPAddress, LogonType

● MASTER

## 8. QUICK QUERY TEMPLATES

### Universal Investigation Template

```
// Copy this skeleton for any investigation
let timeframe = 24h;
let target = 'user@domain.com';

TableName
| where TimeGenerated > ago(timeframe) // 1. Time filter FIRST
| where SomeColumn == target // 2. Narrow scope
| project col1, col2, col3 // 3. Keep useful columns
| summarize count() by col1 // 4. Aggregate
| sort by count_ desc // 5. Rank results
| take 100 // 6. Limit output
```

### Threat Hunting Starter Queries

```
// Who logged in outside business hours? (before 7am or after 8pm)
SignIns
| where TimeGenerated > ago(7d)
| extend HourUTC = datetime_part('hour', todatetime(TimeGenerated))
| where HourUTC < 7 or HourUTC > 20
| where ResultType == '0'
| summarize AfterHoursLogins = count() by UserPrincipalName, HourUTC
| sort by AfterHoursLogins desc

// New user accounts created in last 24h
AuditLogs
| where TimeGenerated > ago(24h)
| where OperationName == 'Add user'
| project TimeGenerated, InitiatedBy, TargetResources

// Service accounts logging in interactively (impossible)
SignIns
| where TimeGenerated > ago(24h)
| where UserPrincipalName startswith 'svc'
| where IsInteractive == true
| project TimeGenerated, UserPrincipalName, IPAddress, AppDisplayName
```

● ADVANCED

## 9. PERFORMANCE & BEST PRACTICES

	Best Practice
■ DO	Put TimeGenerated filter FIRST in every query
■ DO	Use has instead of contains for whole-word text search
■ DO	Use dcount() for approximate distinct counts (faster than count(distinct))
■ DO	Filter early — reduce rows before joining or summarizing
■ DO	Use let statements to make queries readable and reusable

■ DO	Name your summarize columns (e.g. count() as Events, not just count())
■ AVOID	Running queries without a time filter — scans all data
■ AVOID	Using contains when has will do — contains is a full scan
■ AVOID	Joining large tables without filtering both sides first
■ AVOID	Using order by before summarize — sort at the very end
■ AVOID	Smart/curly quotes — always use straight quotes ' or "

## Keyboard Shortcuts in ADX Web UI

Shortcut	Action
Shift + Enter	Run selected query only
F5	Run entire query
Ctrl + Space	Trigger autocomplete
Ctrl + K + C	Comment selected lines
Ctrl + /	Toggle comment
Ctrl + Z	Undo
Ctrl + Shift + F	Format / pretty-print query

KQL Cheat Sheet — Azure Data Explorer · Log Analytics · Microsoft Sentinel | Covers: Beginner → Intermediate → Advanced → Master | Practice dataset:  
SignIns table from generated CSV